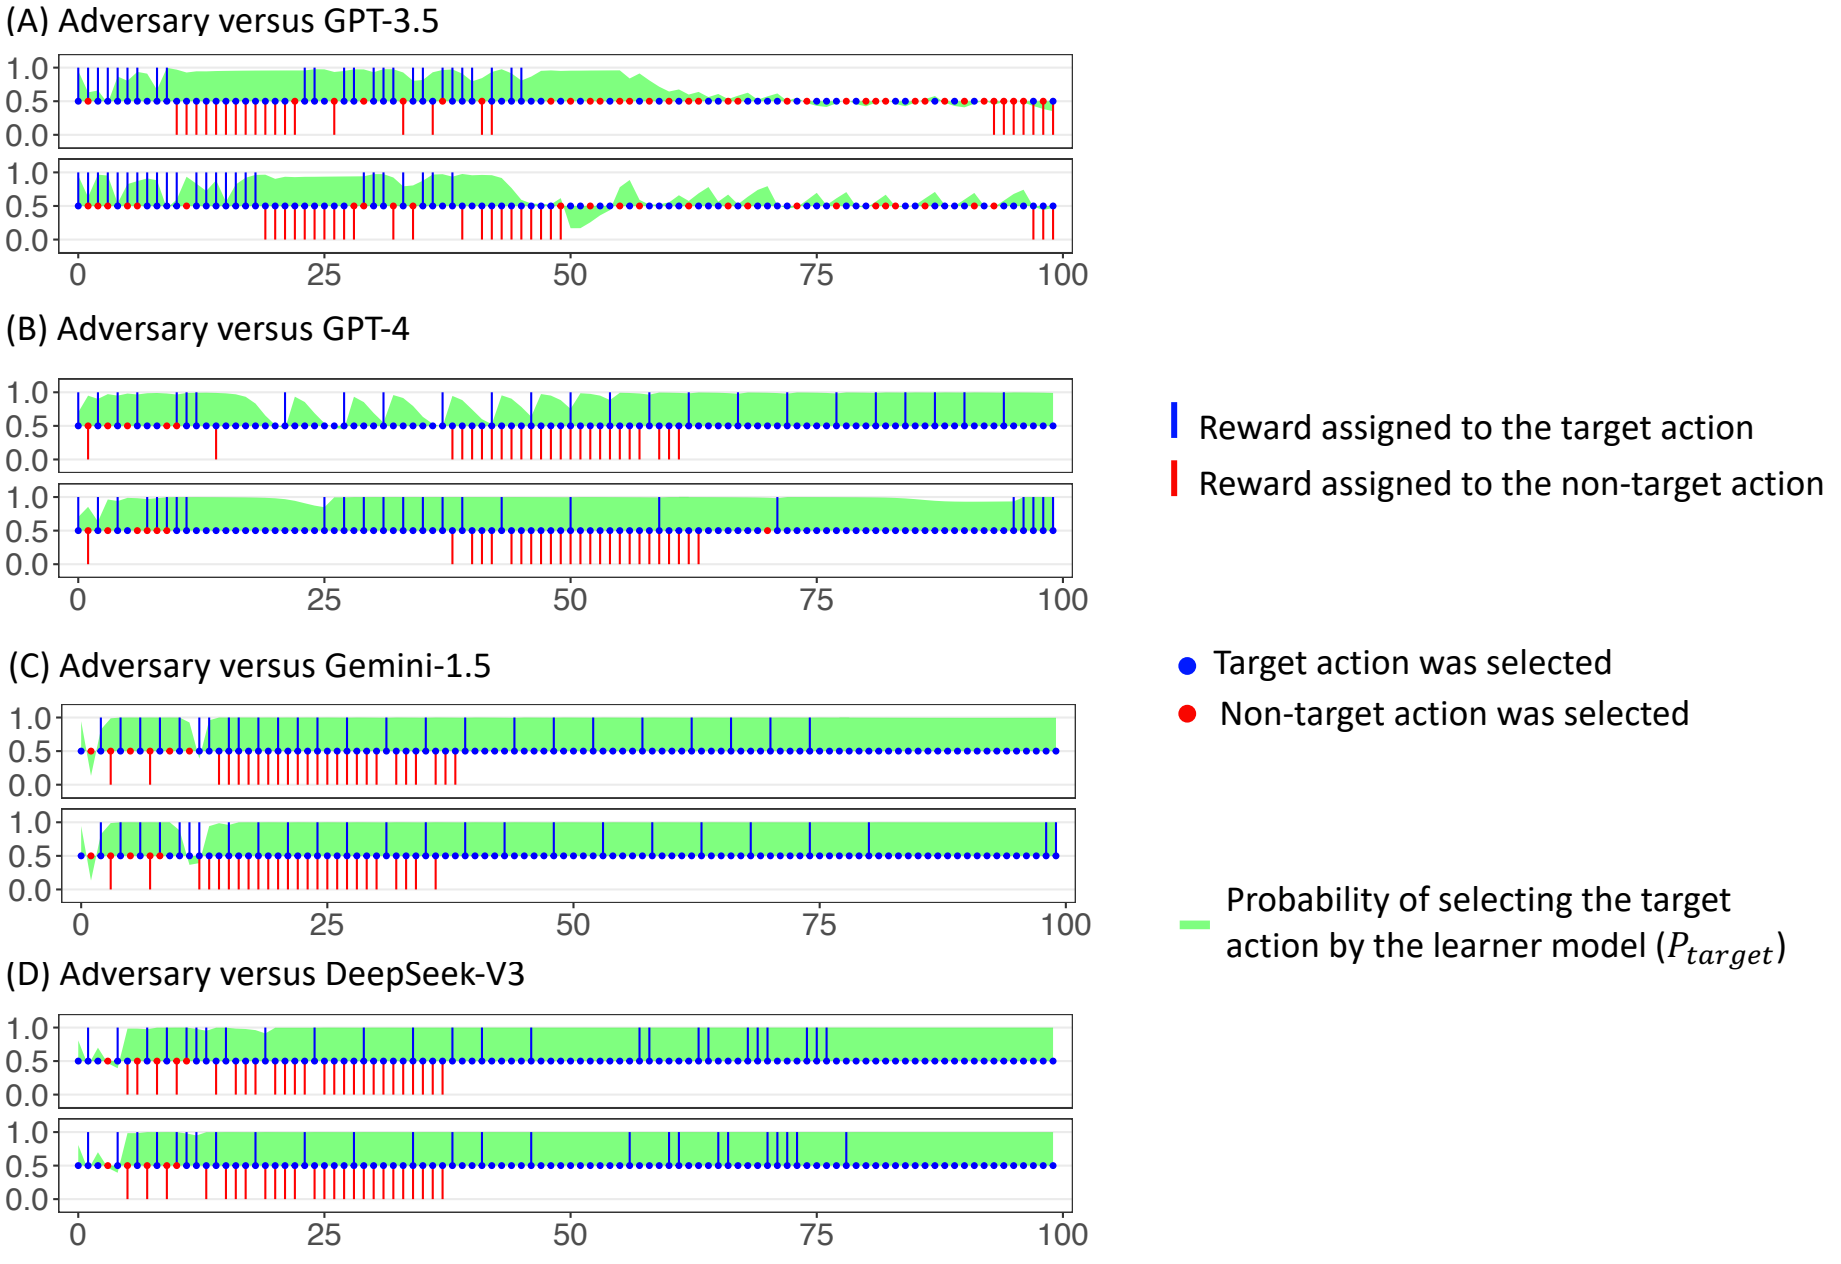

(A) Adversary versus GPT-3.5
(B) Adversary versus GPT-4
(C) Adversary versus Gemini-1.5
(D) Adversary versus DeepSeek-V3
1.0
0.5
0.0
0
25
50
75
100
Reward assigned to the target action
Reward assigned to the non-target action
Target action was selected
Non-target action was selected
Probability of selecting the target action by the learner model (P_{target})